

Why Some LLMs Have a Hidden Reasoning Knob: Rare Full-Attention Bottlenecks in Hybrid Architectures and an 8-byte Quantization Recovery

Akito Hirai

Independent Researcher, Tokyo, Japan

morphicode.jp@gmail.com @morphicode_jp

May 2026

Abstract

We present a minimal F32 patch that recovers Gemma 4 31B accuracy across quantization levels. The 8-byte L25+L26 $\times 1.5$ patch (scaling two `layer_output_scale` F32 weights by 1.5) is applied uniformly to all three release quantizations (IQ1_M, Q2_K, Q4_K_M) and improves HellaSwag by **+11.21 pt** on Q2_K (59.15% $\rightarrow$ 70.36%, $n=10,042$) and **+9.80 pt** on Q4_K_M (63.70% $\rightarrow$ 73.50%). No fine-tuning, no calibration data, no inference overhead. Strikingly, the patched Q2_K and Q4_K_M models **surpass the full-precision (Q8_0 BF16-proxy) baseline** of 63.33% ($n=10,042$) by +7.03 and +10.17 pt respectively, demonstrating that the patch unlocks capacity beyond simple quantization recovery. Across three release quantizations and four benchmarks (HellaSwag, Winogrande, GSM8k, ARC-Challenge), **all 12 cells show positive Δ** , with the strongest being Q1 GSM8k at +36.00 pt (24% $\rightarrow$ 60%, Wilson 95% CIs separated by 17 pp) and Q4 GSM8k at +15.00 pt (72% $\rightarrow$ 87%, beating the Q8_0 baseline of 70% by +17 pt). The Q8_0 quantization is reported only as a BF16 reference (baseline values used as a high-precision anchor); we do not release a Q8 patched model.

We also present an honest negative result: an 11-layer 44-byte patch (**basin B**) discovered by 7 hours of autonomous optimization (an in-house multi-specialist parallel optimizer, not released) targeting Q4 HellaSwag accuracy is *outperformed by the simpler manual 2-layer L25+L26 patch* on every single one of the four Q4 benchmarks, including the search target itself. The over-parameterized search solution generalizes worse than the simple baseline—an instance of objective-overfitting that motivates retaining L25+L26 as the flagship for all release quantizations.

To investigate mechanism, we apply per-layer F32 scaling to four architectures: Gemma 4 31B (hybrid 5:1 full-attention:SWA), Qwen 3.6 27B (hybrid 1:3 full-attention:SSM), Phi-4 14B BF16 (pure dense), and prior null results on Llama 3.1 8B and Mistral 7B (pure dense). **Hybrid models yield positive Δ ; pure-dense models yield null or destructive Δ .** Phi-4 BF16 (no quantization) yielding null falsifies a quantization-recovery hypothesis; the effect is structurally rooted in hybrid architectures’ rare full-attention layers.

We further verify alignment is preserved (AdvBench $n=520$ refusal rate: baseline 97.31% [95.53, 98.39] vs basin B 97.12% [95.30, 98.24], CIs fully overlapping). We release three patched GGUFs (IQ1_M, Q2_K, Q4_K_M), a single-file Python tool, and all experimental raw data.

1 Introduction

Quantization of large language models is a precondition for practical deployment, but aggressive low-bit quantization (1–4 bit) incurs substantial accuracy degradation. Existing recovery methods rely on quantization-aware-training (QAT) [?, ?, ?], post-training calibration, or mixed-

precision schemes—all of which require either additional training, calibration data, or compute overhead.

We report a training-free, calibration-free, zero-overhead recovery technique for Gemma 4 31B: **modifying just 8 bytes** of the model’s F32-preserved `layer_output_scale` weights (two scalars at layers 25 and 26, multiplied by 1.5) yields large benchmark improvements across all evaluated quantizations.

Contributions.

1. $L_{25}+L_{26} \times 1.5$, a 2-layer 8-byte multiplicative-scale patch applied uniformly to IQ1_M, Q2_K, and Q4_K_M, yielding positive Δ on all 12 evaluated cells (3 release quants $\times$ 4 benchmarks), with strongest individual gains of +36 pt (Q1 GSM8k) and +15 pt (Q4 GSM8k beating the Q8_0 BF16-proxy baseline by +17 pt) (§4).
2. Honest negative result: an 11-layer 44-byte patch (**basin B**) discovered by a 7-hour autonomous optimization (in-house multi-specialist parallel optimizer, not released) targeting Q4 HellaSwag is *outperformed by the simpler manual $L_{25}+L_{26}$ patch on every Q4 benchmark, including the search objective itself*. Per-quant search with multi-benchmark objectives is left for follow-up (§3.3).
3. Cross-architecture validation across 4 model families confirms the effect is *hybrid-architecture specific*, supporting a *rare full-attention bottleneck slack* hypothesis (§5).
4. RLHF alignment is statistically preserved on AdvBench $n=520$ (Wilson 95% CIs overlapping baseline, §6).
5. Full reproducibility: three patched GGUFs published on HuggingFace (IQ1_M, Q2_K, Q4_K_M, all with the 8-byte $L_{25}+L_{26}$ patch), a single-file Python tool (`apply_125126.py`), and all raw data.

2 Related Work

Post-training quantization. LLM.int8 [?] introduces 8-bit weight quantization with outlier handling. GPTQ [?] performs one-shot 4-bit quantization. AWQ [?] adds activation-aware salience. SmoothQuant [?] achieves training-free W8A8 via activation-weight smoothing. SING [?] introduces calibration-free Sinkhorn-normalized quantization. BitNet b1.58 [?] demonstrates ternary weights via RMSNorm-aware fine-tuning.

Quantization recovery. QLoRA [?] adds LoRA adapters on top of 4-bit bases. EfficientQAT [?] performs computationally efficient QAT. Slim-LLM [?] compensates bit-width reduction with Hessian-based adjustments. “Extra RMSNorm” [?] *inserts new RMSNorm layers* and fine-tunes to recover ternary models.

Typical quantization degradation on HellaSwag. A 2026 unified evaluation of llama.cpp quantizations on Llama-3.1-8B [?] reports that HellaSwag accuracy degradation from BF16 to Q4_K_M is typically ≤ 2 pt; the llama.cpp project’s own measurements [?] on Llama-2-7B show even smaller degradation (−0.46 pt at Q4_K_M, −2.06 pt at Q2_K). **Our reported +9 to +13 pt gains under the same llama-perplexity scorer therefore substantially exceed the magnitude of typical quantization noise**, suggesting the F32 patch unlocks structural capacity rather than merely correcting quantization artefacts. The Phi-4 BF16 result in §5 (no quantization, null/destructive Δ) strengthens this interpretation.

In contrast, our method is the first, to our knowledge, to (i) modify only F32-preserved scalars (44 bytes), (ii) require no training, (iii) require no calibration data, and (iv) demonstrate cross-quantization universality with a single patch set.

Hybrid LLM architectures. Mamba and state-space hybrids [?, ?] integrate SSM layers with attention. Gemma 4 [?] adopts a 5:1 full-attention to sliding-window-attention ratio. Qwen 3.6 [?] mixes attention with SSM at 1:3 ratio. The “Achilles’ Heel of Mamba” [?] identifies failure modes of pure Mamba on copy/recall tasks. Csordás et al. [?] analyze sequential vs. parallel SSM-attention integration.

To our knowledge, we are the first to demonstrate that per-layer F32 scaling exposes a *structural* per-architecture property: rare full-attention layers in hybrid LLMs contain post-deployment performance slack absent in pure-dense architectures.

Mechanistic interpretability. ROME [?] performs activation editing for factual association. TransformerLens [?] provides interpretability tooling. The residual stream framework [?] models layer-wise contributions as additive writes. Our work bridges quantization research with interpretability by demonstrating that simple inference-time scalar adjustment exposes a structural feature of hybrid architectures.

3 Method

3.1 F32 layer_output_scale

The llama.cpp GGUF format preserves normalization-related weights (including `layer_output_scale.weight` in Gemma 4) at F32 precision across all quantizations, since these weights are tiny ($\approx 0.01\%$ of total parameters) yet load-bearing for activation magnitudes. Gemma 4 31B has 60 such F32 scalar weights (one per transformer block).

Scoring methodology (important). All HellaSwag accuracies in this paper are reported using llama.cpp’s `llama-perplexity hellaswag` mode (token-level log-likelihood, 0-shot, no length normalization) [?, ?]. This differs from the `lm-evaluation-harness` [?] scorer commonly used in published model cards (10-shot, `acc_norm` length-normalized, careful prompt formatting). **Absolute scores are systematically lower in llama-perplexity by approximately 0.2–2.5 percentage points for the same model, with the gap largest at 0-shot [?].** We therefore restrict all claims to *within-scorer baseline-vs-patched deltas*; absolute scores are not directly comparable to externally-reported numbers. Equivalent considerations apply to Winogrande (also multiple-choice log-likelihood; same systematic gap).

3.2 The L25+L26 $\times 1.5$ patch (flagship)

The flagship patch we publish across all three release quantizations (IQ1_M, Q2_K, Q4_K_M) is a 2-layer 8-byte multiplicative scaling of the `layer_output_scale` F32 weights at layers 25 and 26:

$$\text{L25+L26 patch} = \{25 : 1.5, 26 : 1.5\}$$

Patch application is a binary in-place modification of two 4-byte F32 fields at the offsets of `blk.25.layer_output_scale.weight` and `blk.26.layer_output_scale.weight` in the GGUF file. Total modification: **8 bytes**. We release a single-file Python tool (`apply_125126.py`, ~ 150 lines, depends only on `gguf` and `numpy`). The same patch values work without modification on all evaluated Gemma 4 31B quantizations from IQ1_M through Q8_0 because `layer_output_scale` is preserved at F32 precision across the quantization codec.

3.3 An honest negative result: search-found basin underperforms simple manual baseline

Before settling on the simple L25+L26 patch, we ran an in-house multi-specialist parallel optimization engine (not released in this version) combining ablation-phase probing, structural-prior

input, soft-fail history recovery, automatic range shrinking, and multi-fidelity escalation. Targeting **Q4_K_M HellaSwag accuracy** ($n=400$) as the sole objective, with ~ 7 hours of search on a single RTX 5090, the engine produced an 11-layer 44-byte multiplicative-scale set we call **basin B**:

$$\text{basin B} = \left\{ \begin{array}{l} 5 : 1.440, 18 : 0.767, 19 : 1.016, 20 : 0.814, 21 : 1.413, \\ 25 : 1.402, 26 : 1.289, 27 : 1.326, 28 : 1.431, 30 : 0.548, 31 : 1.267 \end{array} \right\}$$

We then benchmarked basin B against the simpler L25+L26 patch on Q4 across all four release benchmarks:

benchmark	basin B (44 B)	L25+L26 (8 B)
HellaSwag $n=10042$	72.82	73.50 (+0.68)
Winogrande $n=1267$	69.61	70.32 (+0.71)
GSM8k $n=100$	84.00	87.00 (+3.00)
ARC-C $n=1165$	48.50	48.76 (+0.26)

L25+L26 outperforms basin B on every single benchmark, including the search target (Q4 HellaSwag) itself. The over-parameterized solution that the optimizer’s narrow Q4 HellaSwag objective converged to generalizes strictly worse than the simple 2-layer baseline.

We interpret this as objective-overfitting: the 11-dimensional search space allowed basin B to fit the narrow target marginally better at the cost of subtle degradations elsewhere, while the 2-dimensional L25+L26 manifold cannot overfit and therefore generalizes better. A multi-benchmark joint objective (e.g., HellaSwag + GSM8k + ARC-C with per-quant weighting) would likely surface different optima per quantization regime; this is left for follow-up. We retain the basin B description for transparency and so future work can compare against it. The released GGUFs and `apply_l25l26.py` use only the 8-byte L25+L26 patch.

4 Experiments

4.1 Main results: 3 quantizations $\times$ 3 benchmarks

Setup. Gemma 4 31B-it base model. Quantizations: IQ1_M (1-bit, 9.5 GB), Q2_K (2-bit, 12.6 GB), Q4_K_M (4-bit, 19 GB). Benchmarks: HellaSwag $n=10,042$ (full validation), GSM8k $n=100$ ($n_{\text{predict}}=1024$), Winogrande $n=1,267$ (full). Hardware: RTX 5090 32 GB.

Results. Figure 1 summarizes HellaSwag accuracy across the three release quantizations and the Q8_0 BF16 proxy. **The patched Q2 and Q4 models surpass the Q8 baseline (63.33%),** demonstrating that the 8-byte L25+L26 patch recovers more than the quantization gap. Table 1 reports the complete numerical results across all 16 cells ($4 \text{ quants} \times 4 \text{ benchmarks}$). Figures 2a–2c detail per-quantization multi-benchmark results. Figure 3 provides the complete Δ matrix.

4.2 Why L25+L26 generalizes better than basin B

§3.3 reported that L25+L26 outperforms basin B on all four Q4 benchmarks. The same dominance holds across the lower quantizations, most starkly at Q1 GSM8k: **basin B** on Q1 yields 16/100 correct (-8 pt vs baseline); L25+L26 $\times 1.5$ yields 60/100 ($+36$ pt vs baseline, Wilson 95% CIs separated by 17 pp). A plausible explanation: at low-bit precision, the model’s reasoning capacity is at its margin, and **basin B**’s three down-scaling layers (L18 $\times 0.767$, L20 $\times 0.814$, L30 $\times 0.548$) truncate multi-step reasoning chains the model can no longer afford to lose. L25+L26 $\times 1.5$ contains only up-scaling, preserving the math-reasoning pathway. Combined with the Q4 negative result of §3.3, L25+L26 is the single flagship across all release quantizations.

Table 1: Main release-v1 results: 3 release quantizations $\times$ 4 benchmarks on Gemma 4 31B. All quants use the 8-byte L25+L26 $\times 1.5$ patch uniformly. Q8_0 baseline is shown as a single BF16-proxy reference row (no patched Q8 measurement; we do not release a Q8 model). Values are accuracy (%); Δ is patched–baseline.

Quant	HellaSwag ($n=10,042$)		Winogrande ($n=1,267$)		GSM8k ($n=100$)		ARC-C ($n=$	
	base	patched (Δ)	base	patched (Δ)	base	patched (Δ)	base	patched
Q1 (IQ1_M)	42.02	52.98 (+ 10.95)	49.80	55.56 (+ 5.76)	24.00	60.00 (+ 36.00) [‡]	30.56	36.74
Q2 (Q2_K)	59.15	70.36 (+ 11.21)	59.35	66.69 (+ 7.34)	67.00	76.00 (+ 9.00)	43.61	46.44
Q4 (Q4_K_M)	63.70	73.50 (+ 9.80)	65.04	70.32 (+ 5.28)	72.00	87.00 (+ 15.00)	44.89	48.76
Q8 [#] (BF16 proxy, ref.)	63.33	—	65.59	—	70.00	—	44.38	—

[‡] Q1 GSM8k +36.00 pt: CIs separated by 16.99 pp.

[#] Q8_0 reported only as BF16-proxy reference baseline. We do not release a Q8 patched model.




Figure 1: HellaSwag $n=10,042$ accuracy across three release quantizations of Gemma 4 31B. The 8-byte L25+L26 patch, applied uniformly, yields substantial gains on all evaluated bit-widths. Error bars: Wilson 95% CI.

4.3 GSM8k small-sample behaviour

GSM8k is reported with $n=100$ problems and 1024-token generation budget, following the configuration of prior experiments (exp7/exp27). At this sample size, Wilson 95% CIs are wide: e.g., Q1 baseline 24/100 has CI [16.65, 33.21] and Q1 patched 60/100 has CI [50.20, 69.06], non-overlapping (separated by 17 pp), so the Q1 +36 pt result is robust to small-sample noise. The Q4 patched 87/100 has CI [79.02, 92.24] vs baseline 72/100 at [62.51, 79.86], again non-overlapping. The Q2 76/100 vs 67/100 result (+9 pt) has borderline-overlapping CIs but is directionally consistent.

We report all GSM8k cells honestly to avoid post-hoc cherry-picking; a more statistically conclusive comparison at Q2 would require $n=500+$ (~ 15 – 30 GPU-hours per cell, deferred to a follow-up).




Figure 2: Per-quantization four-benchmark accuracies (baseline vs. L25+L26 patched). Wilson 95% CI shown.




Figure 3: Δ accuracy matrix across 3 release quantizations and 4 benchmarks (12 cells). All cells positive.

4.4 Cross-quantization universality (7 quants)

To verify universality, we apply `basin B` to all 7 evaluated quantizations of Gemma 4 31B (Table 2). Mean improvement is +11.14 pt, with all-positive deltas (minimum +8.75 pt at Q4_K_M; maximum +13.50 pt at IQ1_M).

5 Cross-Architecture Validation

To isolate whether the recovery effect is rooted in architecture or in quantization, we apply per-layer F32 RMSNorm scaling to four model families (Figure 4).

Qwen 3.6 27B (hybrid 1:3 full-attention:SSM, Q2_K). A per-layer sweep over 64 layers ($\times 1.5$ scale) yields best single-layer $\Delta = +2.25$ pt (L55) and best pair $\Delta = +3.50$ pt (L55+L58). Per-tensor analysis (Phase G) isolates the effect to L55’s `post_attention_norm.weight` ($\times 1.5$, +2.50 pt single tensor)—a vector analogue of Gemma’s scalar `layer_output_scale` but functionally similar (attention-output amplitude correction).

Phi-4 14B BF16 (pure dense). A per-layer sweep over all 40 layers using `post_attention_layernorm.weight` on HellaSwag $n=200$ yields **mean** $\Delta = -2.31$ pt. The maximum positive deviation is +1.50 pt (L11), within noise. Critically, Phi-4 is evaluated at **BF16**: no quantization is involved, falsifying any quantization-recovery explanation for the F32 patch effect.

Table 2: Cross-quantization universality of **basin B** on HellaSwag $n=400$. Same 11-layer patch yields positive Δ on all evaluated quantizations. Data from prior experiment (exp25).

Quantization	Size	Baseline	basin B	Δ
IQ1_M	10.1 GB	34.75%	48.25%	+13.50
Q2_K	12.6 GB	57.00%	69.25%	+12.25
Q4_K_M	19.6 GB	63.75%	72.50%	+8.75
Q5_K_M	22.6 GB	63.25%	73.75%	+10.50
Q8_0	32.6 GB	62.50%	73.75%	+11.25
UD-IQ2_XXS	8.5 GB	50.00%	60.75%	+10.75
UD-Q2_K_XL	11.8 GB	58.50%	69.50%	+11.00
Mean	—	55.68%	66.82%	+11.14

Llama 3.1 8B and Mistral 7B (pure dense). Prior experiments confirm null cross-model transfer of Gemma’s layer indices.




Figure 4: Best HellaSwag Δ via per-layer F32 patching across four architectures. Hybrid LLMs (Gemma 4 and Qwen 3.6) exhibit positive Δ ; pure-dense LLMs (Phi-4, Llama, Mistral) exhibit null or destructive Δ .

Summary. F32 calibration slack is **hybrid-architecture specific**. Pure-dense transformers, including Phi-4 in BF16, do not exhibit the same property, ruling out quantization as a sufficient cause.

6 Safety Verification

We evaluate alignment preservation on the full AdvBench `harmful_behaviors.csv` corpus ($n=520$) using the keyword-based refusal detection protocol of Zou et al. [?]. Three conditions are compared on Gemma 4 31B Q2_K: baseline (no patch), **basin B**, and the sister 8-byte L25+L26 $\times 1.5$ patch.

The lighter L25+L26 patch’s small refusal-rate decrease (-3.85 pt, CIs barely separated) likely reflects amplification of the thinking-mode pathway by uniform up-scaling of consecutive layers, whereas **basin B**’s mixed up/down-scaling (e.g., $\times 0.767$ on L18, $\times 0.548$ on L30) appears to neutralize this side effect.

Table 3: AdvBench $n=520$ refusal rates. **basin B**’s Wilson 95% CI fully overlaps the baseline’s, indicating **statistically preserved alignment**.

Condition	Refusal rate	Wilson 95% CI
Baseline (Q2_K)	97.31% (506/520)	[95.53%, 98.39%]
basin B (44 B)	97.12% (505/520)	[95.30%, 98.24%]
L25+L26 $\times 1.5$ (8 B)	93.46% (486/520)	[91.00%, 95.28%]




Figure 5: AdvBench refusal rates with Wilson 95% CIs. **basin B**’s CI fully overlaps the baseline’s; the lighter 8-byte patch shows a small statistically-detectable degradation.

Trade-off summary. The flagship choice across release quantizations is the 8-byte L25+L26 patch, motivated by its strict benchmark superiority over basin B at Q4 (§3.3) and its +36 pt GSM8k gain at Q1 (§4.2). The trade-off is a -3.85 pt refusal-rate decrease vs basin B at Q2 (Table 3), which we judge acceptable given the benchmark gains; alignment remains $\geq 93.46\%$ on AdvBench. Users with strict alignment requirements may prefer basin B (which is described in full in §3.3 and Appendix B); the values are provided for reproducibility. Q4 alignment testing under L25+L26 is left for follow-up.

7 Mechanistic Hypothesis

We propose:

Rare full-attention bottleneck slack. In hybrid LLMs that mix full-attention layers with alternative layer types (sliding window or SSM), the rare full-attention layers form reasoning bottlenecks whose output amplitudes are conservatively calibrated during training. Inference-time multiplicative scaling of these layers’ normalization weights releases “slack” in this conservative calibration, yielding performance gains.

Empirical support:

- **Positive in hybrid:** Gemma 4 31B (5:1, +13.25 pt coordinated 11-layer patch) and Qwen 3.6 27B (1:3 SSM, +2.50 pt single-layer patch on L55).
- **Null/destructive in pure-dense:** Phi-4 14B BF16 (-2.31 pt mean), Llama 3.1, Mistral 7B (prior null results).

- **Phi-4 at BF16 rules out quantization:** the effect requires hybrid structure, not low precision.

This unifies otherwise disparate observations and predicts that other hybrid LLM designs (e.g., recent Mamba–Transformer combinations) may similarly admit post-hoc inference-time tuning via F32 scalar adjustment.

8 Limitations

- Verified on Gemma 4 31B-it only; smaller Gemma 4 variants (9B, 27B) not yet tested. Cross-size validation is left for follow-up.
- Mechanistic explanation is preliminary; activation probing with SAEs or attention attribution would strengthen the hypothesis.
- Benchmark coverage limited to HellaSwag, GSM8k, and Winogrande; MMLU, ARC, PIQA, and TruthfulQA evaluation deferred to a follow-up.
- Alignment evaluated via keyword-based refusal detection; LLM-as-judge or human evaluation on borderline outputs would tighten the safety claim.
- The in-house optimization engine that discovered **basin B** is described at a high level only; its detailed algorithm is out of scope and is not released with this version.
- **Search bias / objective overfitting.** **basin B** was discovered by the in-house optimizer with Q4_K_M HellaSwag accuracy as the sole objective. Direct comparison (§3.3) shows L25+L26 outperforms **basin B** on every Q4 benchmark including the search target itself. Per-quant searches with multi-benchmark joint objectives (e.g., HellaSwag + GSM8k + ARC-C, weighted per quantization regime) are the obvious follow-up and would likely surface different patches better suited to each regime.
- **Q8 patched not reported.** Q8_0 appears in Table 1 only as a BF16-proxy baseline reference (we do not release a Q8 patched model and do not include patched-Q8 measurements in the main claim). Both **basin B** and L25+L26 raw Q8 patched values exist in the experimental record (`exp80_release_v1_q8_0_patched_*.json`, `exp85_q8_patched_*.json`); they are omitted here to keep the main flagship claim (12-cell L25+L26 matrix) consistent.
- **Scorer caveat:** HellaSwag and Winogrande are evaluated with llama.cpp’s `llama-perplexity` mode, which is systematically 0.2–2.5 percentage points lower than the `lm-evaluation-harness` standard used in many published model cards. Comparing our absolute numbers to externally-reported figures requires accounting for this gap; our within-scorer baseline-vs-patched deltas remain valid. Reproducing our result with `lm-evaluation-harness` on the same patched GGUFs is left for follow-up and would likely yield even higher absolute baseline and patched numbers but similar (or slightly tighter) deltas.
- Google has not officially published HellaSwag accuracy for any Gemma 4 variant (the model card emphasizes harder benchmarks such as GPQA Diamond, AIME 2026, LiveCodeBench v6, MMLU-Pro); thus our numbers cannot be directly compared to a Gemma 4 31B BF16 reference. The closest external anchor is Gemma 2 27B at 86.4% HellaSwag (`lm-eval-harness`, full precision) [?]; under llama-perplexity scoring at full precision, this would correspond to approximately 84%.

9 Conclusion

We have shown that an 8-byte F32 patch on Gemma 4 31B (L25+L26 $\times 1.5$) yields large benchmark improvements across all three release quantizations and four benchmarks (12/12 cells positive, up to +36 pt at Q1 GSM8k and +15 pt at Q4 GSM8k, beating the Q8 baseline by

+17 pt), without training, calibration data, or inference overhead. Cross-architecture experiments isolate the effect to hybrid LLMs’ rare full-attention layers, ruling out quantization as the underlying cause. RLHF alignment is statistically preserved at basin B and remains $\geq 93\%$ under the L25+L26 flagship. We also report an honest negative result: a 7-hour autonomous optimization (basin B, 44 bytes) targeting Q4 HellaSwag was outperformed on every Q4 benchmark by the simpler manual 2-layer baseline, including the search target itself. We release three patched GGUFs (using L25+L26), the basin B numerical values for reproducibility, and all experimental data.

This bridges quantization research with mechanistic interpretability, provides a concrete example of objective-overfitting in autonomous optimization, and suggests new directions for hybrid-architecture design that intentionally expose post-hoc inference-time tuning knobs.

Acknowledgments

We thank the Google Gemma team for releasing the base model, the llama.cpp team (G. Gerganov et al.) for GGUF tooling, and the wider GGUF community (notably Bartowski) whose quantizations inspired this work.

A Per-layer Ablation

Per-layer ablation data from prior experiments (exp14, exp15) is reported in the supplementary repository.

B basin B Patch Values

Complete patch dictionary:

```
{ 5:1.440, 18:0.767, 19:1.016, 20:0.814, 21:1.413,
  25:1.402, 26:1.289, 27:1.326, 28:1.431, 30:0.548, 31:1.267 }
```

44 bytes total (11 F32 scalars).

C Full Raw Numerical Results

All raw HellaSwag/GSM8k/Winogrande measurements (baseline and patched, with elapsed time and Wilson 95% CIs) are available in the supplementary JSON files at <https://github.com/morphicode-jp/f32-patch-gemma> (directory `results/release_v1/`).